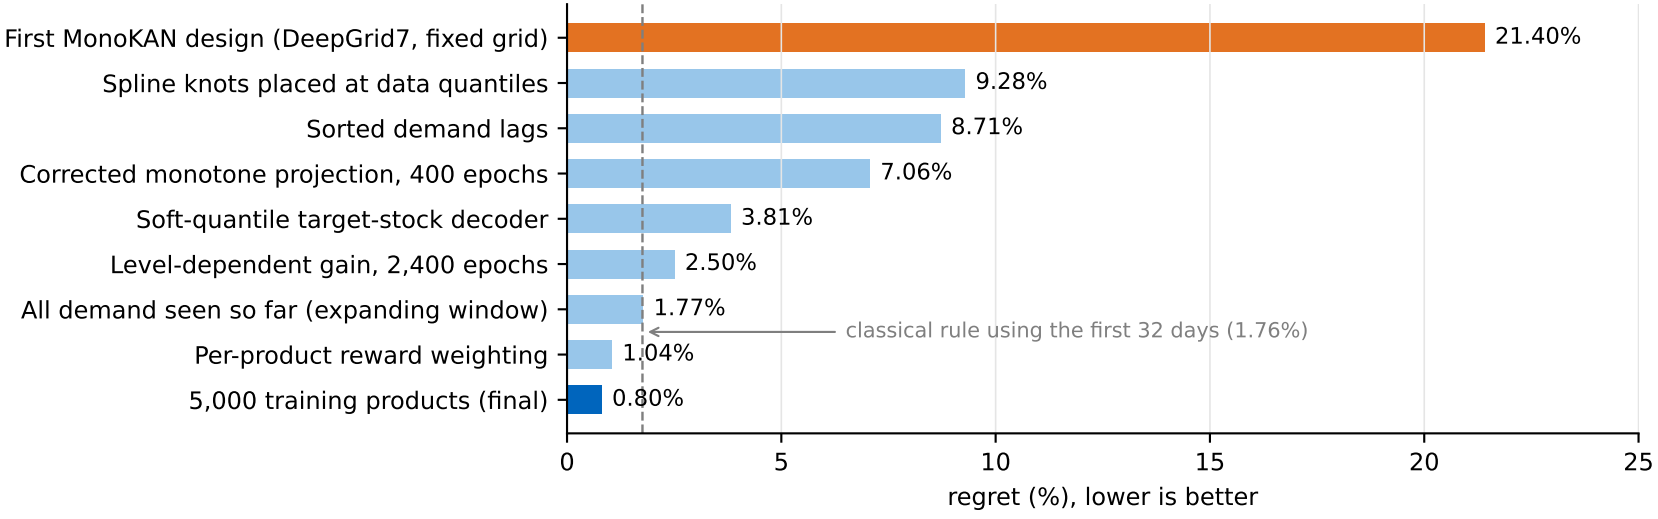
First MonoKAN design (DeepGrid7, fixed grid)
Spline knots placed at data quantiles
Sorted demand lags
Corrected monotone projection, 400 epochs
Soft-quantile target-stock decoder
Level-dependent gain, 2,400 epochs
All demand seen so far (expanding window)
Per-product reward weighting
5,000 training products (final)
21.40%
9.28%
8.71%
7.06%
3.81%
2.50%
1.77%
1.04%
0.80%
classical rule using the first 32 days (1.76%)
0
5
10
15
20
25
regret (%), lower is better